

sms_spam_detection

August 23, 2025

```
[1]: import numpy as np
import pandas as pd
```

```
[2]: import sys
print(sys.executable)
```

/Users/ayush/Desktop/spam-classifier/myenv/bin/python

```
[3]: df=pd.read_csv('spam.csv', encoding='latin1')
df.head(5)
```

```
[3]:      v1                                                    v2 Unnamed: 2  \
0   ham  Go until jurong point, crazy.. Available only ...      NaN
1   ham                                Ok lar... Joking wif u oni...      NaN
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...      NaN
3   ham  U dun say so early hor... U c already then say...      NaN
4   ham  Nah I don't think he goes to usf, he lives aro...      NaN

      Unnamed: 3 Unnamed: 4
0           NaN          NaN
1           NaN          NaN
2           NaN          NaN
3           NaN          NaN
4           NaN          NaN
```

```
[4]: df.shape
```

```
[4]: (5572, 5)
```

1.Data Cleaning 2.EDA 3.Text Preprocessing 4.Model Building 5.Evaluation 6.Improvement 7.Web-site 8.Deployment

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0    v1          5572 non-null   object
```

```

1   v2          5572 non-null   object
2   Unnamed: 2    50 non-null   object
3   Unnamed: 3    12 non-null   object
4   Unnamed: 4     6 non-null   object
dtypes: object(5)
memory usage: 217.8+ KB

```

```
[6]: df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], inplace=True)
```

```
[7]: df.sample()
```

```
[7]:      v1          v2
429  ham  I wnt to buy a BMW car urgently..its vry urgen...
```

```
[8]: df.rename(columns={'v1': 'target', 'v2': 'text'}, inplace=True)
```

```
[9]: from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
```

```
[10]: df['target']=encoder.fit_transform(df['target'])
```

```
[11]: df.head(5)
```

```
[11]:   target          text
0      0  Go until jurong point, crazy.. Available only ...
1      0              Ok lar... Joking wif u oni...
2      1  Free entry in 2 a wkly comp to win FA Cup fina...
3      0  U dun say so early hor... U c already then say...
4      0  Nah I don't think he goes to usf, he lives aro...
```

```
[12]: df.isnull().sum()
```

```
[12]: target      0
text          0
dtype: int64
```

```
[13]: df.duplicated().sum()
```

```
[13]: np.int64(403)
```

```
[14]: df=df.drop_duplicates(keep='first')
```

```
[15]: df.duplicated().sum()
```

```
[15]: np.int64(0)
```

```
[16]: df.shape
```

```
[16]: (5169, 2)
```


//EDA

```
[17]: df['target'].value_counts()
```

```
[17]: target
0      4516
1       653
Name: count, dtype: int64
```

```
[18]: import matplotlib.pyplot as plt
import seaborn as sns

plt.pie(df['target'].value_counts(), labels=['ham', 'spam'], autopct="%0.2f")
plt.show()
```



```
[19]: import nltk
```

```
[20]: nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /Users/ayush/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

```
[20]: True
```

```
[21]: nltk.download('punkt_tab')
```

```
[nltk_data] Downloading package punkt_tab to /Users/ayush/nltk_data...  
[nltk_data]   Package punkt_tab is already up-to-date!
```

```
[21]: True
```

```
[22]: df['num_characters']=df['text'].apply(len)
```

```
[23]: df.head()
```

```
[23]:
```

	target	text	num_characters
0	0	Go until jurong point, crazy.. Available only ...	111
1	0	Ok lar... Joking wif u oni...	29
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	0	U dun say so early hor... U c already then say...	49
4	0	Nah I don't think he goes to usf, he lives aro...	61

```
[24]: df['num_words'] = df['text'].apply(lambda x:len(nltk.word_tokenize(x)))  
df.head()
```

```
[24]:
```

	target	text	num_characters	\
0	0	Go until jurong point, crazy.. Available only ...	111	
1	0	Ok lar... Joking wif u oni...	29	
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	
3	0	U dun say so early hor... U c already then say...	49	
4	0	Nah I don't think he goes to usf, he lives aro...	61	

	num_words
0	24
1	8
2	37
3	13
4	15

```
[25]: df['num_sentences']=df['text'].apply(lambda x:len(nltk.sent_tokenize(x)))  
df.head()
```

```
[25]:
```

	target	text	num_characters	\
0	0	Go until jurong point, crazy.. Available only ...	111	
1	0	Ok lar... Joking wif u oni...	29	
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	
3	0	U dun say so early hor... U c already then say...	49	
4	0	Nah I don't think he goes to usf, he lives aro...	61	

	num_words	num_sentences
0	24	2
1	8	2

2	37	2
3	13	1
4	15	1

```
[26]: df[['num_characters', 'num_words', 'num_sentences']].describe()
```

```
[26]:
```

	num_characters	num_words	num_sentences
count	5169.000000	5169.000000	5169.000000
mean	78.977945	18.455794	1.965564
std	58.236293	13.324758	1.448541
min	2.000000	1.000000	1.000000
25%	36.000000	9.000000	1.000000
50%	60.000000	15.000000	1.000000
75%	117.000000	26.000000	2.000000
max	910.000000	220.000000	38.000000

```
[27]: df[df['target']==0][['num_characters', 'num_words', 'num_sentences']].describe()
```

```
[27]:
```

	num_characters	num_words	num_sentences
count	4516.000000	4516.000000	4516.000000
mean	70.459256	17.123782	1.820195
std	56.358207	13.493970	1.383657
min	2.000000	1.000000	1.000000
25%	34.000000	8.000000	1.000000
50%	52.000000	13.000000	1.000000
75%	90.000000	22.000000	2.000000
max	910.000000	220.000000	38.000000


```
[28]: df[df['target']==1][['num_characters', 'num_words', 'num_sentences']].describe()
```

```
[28]:
```

	num_characters	num_words	num_sentences
count	653.000000	653.000000	653.000000
mean	137.891271	27.667688	2.970904
std	30.137753	7.008418	1.488425
min	13.000000	2.000000	1.000000
25%	132.000000	25.000000	2.000000
50%	149.000000	29.000000	3.000000
75%	157.000000	32.000000	4.000000
max	224.000000	46.000000	9.000000


```
[29]: sns.histplot(df[df['target']==0]['num_characters'])
sns.histplot(df[df['target']==1]['num_characters'], color='red')
```

```
[29]: <Axes: xlabel='num_characters', ylabel='Count'>
```




```
[30]: sns.histplot(df[df['target']==0]['num_words'])  
sns.histplot(df[df['target']==1]['num_words'],color='red')
```

```
[30]: <Axes: xlabel='num_words', ylabel='Count'>
```



```
[31]: sns.pairplot(df,hue='target')
```

```
[31]: <seaborn.axisgrid.PairGrid at 0x14361ef10>
```



```
[32]: sns.heatmap(df.select_dtypes(include=['int64', 'float64']).corr(),annot=True)
```

```
[32]: <Axes: >
```




```
[33]: import nltk
      nltk.download('stopwords')
```

[nltk_data] Downloading package stopwords to /Users/ayush/nltk_data..
[nltk_data] Package stopwords is already up-to-date!

[33]: True

```
[34]: from nltk.corpus import stopwords
      stopwords.words('english')
```

[34]: ['a',
 'about',
 'above',
 'after',
 'again',
 'against',
 'ain',
 'all',
 'am',
 'an',
 'and',
 'any',

'are',
'aren',
"aren't",
'as',
'at',
'be',
'because',
'been',
'before',
'being',
'below',
'between',
'both',
'but',
'by',
'can',
'couldn',
"couldn't",
'd',
'did',
'didn',
"didn't",
'do',
'does',
'doesn',
"doesn't",
'doing',
'don',
"don't",
'down',
'during',
'each',
'few',
'for',
'from',
'further',
'had',
'hadn',
"hadn't",
'has',
'hasn',
"hasn't",
'have',
'haven',
"haven't",
'having',
'he',

"he'd",
"he'll",
'her',
'here',
'hers',
'herself',
"he's",
'him',
'himself',
'his',
'how',
'i',
"i'd",
'if',
"i'll",
"i'm",
'in',
'into',
'is',
'isn',
"isn't",
'it',
"it'd",
"it'll",
"it's",
'its',
'itself',
"i've",
'just',
'll',
'm',
'ma',
'me',
'mightn',
"mightn't",
'more',
'most',
'mustn',
"mustn't",
'my',
'myself',
'needn',
"needn't",
'no',
'nor',
'not',
'now',

'o',
'of',
'off',
'on',
'once',
'only',
'or',
'other',
'our',
'ours',
'ourselves',
'out',
'over',
'own',
're',
's',
'same',
'shan',
"shan't",
'she',
"she'd",
"she'll",
"she's",
'should',
'shouldn',
"shouldn't",
"should've",
'so',
'some',
'such',
't',
'than',
'that',
"that'll",
'the',
'their',
'theirs',
'them',
'themselves',
'then',
'there',
'these',
'they',
"they'd",
"they'll",
"they're",
"they've",

'this',
'those',
'through',
'to',
'too',
'under',
'until',
'up',
've',
'very',
'was',
'wasn',
"wasn't",
'we',
"we'd",
"we'll",
"we're",
'were',
'weren',
"weren't",
"we've",
'what',
'when',
'where',
'which',
'while',
'who',
'whom',
'why',
'will',
'with',
'won',
"won't",
'wouldn',
"wouldn't",
'y',
'you',
"you'd",
"you'll",
'your',
"you're",
'yours',
'yourself',
'yourselves',
"you've"]

```
[35]: import string
      string.punctuation
```

```
[35]: '!"#$%&\'()*+,-./:;<=>?@[\\]^_`{|}~'
```

```
[36]: from nltk.stem.porter import PorterStemmer
      ps = PorterStemmer()
      print(ps.stem("running"))
```

run

```
[37]: def transform_text(text):
      text=text.lower()
      text=nltk.word_tokenize(text)
      y=[]
      for i in text:
          if i.isalnum():
              y.append(i)

      text=y[:]
      y.clear()

      for i in text:
          if i not in stopwords.words('english') and i not in string.punctuation:
              y.append(i)

      text=y[:]
      y.clear()

      for i in text:
          y.append(ps.stem(i))

      return " ".join(y)
```

```
[38]: transform_text("I'm gonna be home soon and i don't want to talk about this_
      ↪stuff anymore tonight, k? I've cried enough today.")
```

```
[38]: 'gon na home soon want talk stuff anymor tonight k cri enough today'
```

```
[39]: df['text'][0]
```

```
[39]: 'Go until jurong point, crazy.. Available only in bugis n great world la e
      buffet... Cine there got amore wat...'
```

```
[40]: df['transformed_text']=df['text'].apply(transform_text)
```

```
[41]: df.head()
```

```
[41]:
```

	target	text	num_characters	\
0	0	Go until jurong point, crazy.. Available only ...	111	
1	0	Ok lar... Joking wif u oni...	29	
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	
3	0	U dun say so early hor... U c already then say...	49	
4	0	Nah I don't think he goes to usf, he lives aro...	61	

	num_words	num_sentences	transformed_text
0	24	2	go jurong point crazi avail bugi n great world...
1	8	2	ok lar joke wif u oni
2	37	2	free entri 2 wkli comp win fa cup final tkt 21...
3	13	1	u dun say earli hor u c already say
4	15	1	nah think goe usf live around though

```
[42]: from wordcloud import WordCloud
wc = WordCloud(width=500,height=500,min_font_size=10,background_color='white')
```

```
[43]: spam_wc=wc.generate(df[df['target']==1]['transformed_text'].str.cat(sep=" "))
plt.figure(figsize=(16,8))
plt.imshow(spam_wc)
```

```
[43]: <matplotlib.image.AxesImage at 0x110a74310>
```



```
[44]: ham_wc=wc.generate(df[df['target']==0]['transformed_text'].str.cat(sep=" "))
plt.figure(figsize=(16,8))
plt.imshow(ham_wc)
```

```
[44]: <matplotlib.image.AxesImage at 0x1100121d0>
```





```
[45]: spam_corpus=[]
      for msg in df[df['target']==1]['transformed_text'].tolist():
          for word in msg.split():
              spam_corpus.append(word)
```

```
[46]: len(spam_corpus)
```

[46] : 9939

```
[47]: from collections import Counter
sns.barplot(x=pd.DataFrame(Counter(spam_corpus).most_common(30))[0], y=pd.
        DataFrame(Counter(spam_corpus).most_common(30))[1])
plt.xticks(rotation='vertical')
```

```
plt.show()
```




```
[48]: ham_corpus=[]  
for msg in df[df['target']==0]['transformed_text'].tolist():  
    for word in msg.split():  
        ham_corpus.append(word)
```

```
[49]: len(ham_corpus)
```

```
[49]: 35404
```

```
[50]: from collections import Counter  
sns.barplot(x=pd.DataFrame(Counter(ham_corpus).most_common(30))[0], y=pd.  
    ↪ DataFrame(Counter(ham_corpus).most_common(30))[1])  
plt.xticks(rotation='vertical')  
plt.show()
```



```
[51]: from sklearn.feature_extraction.text import CountVectorizer,TfidfVectorizer
      cv=CountVectorizer()
      tfidf=TfidfVectorizer(max_features=3000)
```

```
[52]: X=tfidf.fit_transform(df['transformed_text']).toarray()
```

```
[53]: X.shape
```

```
[53]: (5169, 3000)
```

```
[54]: y = df['target'].values
```

```
[55]: from sklearn.model_selection import train_test_split
```

```
[56]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
      ↪2,random_state=2)
```

```
[57]: from sklearn.naive_bayes import GaussianNB,MultinomialNB,BernoulliNB
      from sklearn.metrics import accuracy_score,confusion_matrix,precision_score
```

```
[58]: gnb=GaussianNB()  
      mnb=MultinomialNB()  
      bnb=BernoulliNB()
```

```
[59]: gnb.fit(X_train,y_train)  
      y_pred1=gnb.predict(X_test)  
      print(accuracy_score(y_test,y_pred1))  
      print(confusion_matrix(y_test,y_pred1))  
      print(precision_score(y_test,y_pred1))
```

```
0.8694390715667312  
[[788 108]  
 [ 27 111]]  
0.5068493150684932
```

```
[60]: mnb.fit(X_train,y_train)  
      y_pred2=mnb.predict(X_test)  
      print(accuracy_score(y_test,y_pred2))  
      print(confusion_matrix(y_test,y_pred2))  
      print(precision_score(y_test,y_pred2))
```

```
0.9709864603481625  
[[896   0]  
 [ 30 108]]  
1.0
```

```
[61]: bnb.fit(X_train,y_train)  
      y_pred3=bnb.predict(X_test)  
      print(accuracy_score(y_test,y_pred3))  
      print(confusion_matrix(y_test,y_pred3))  
      print(precision_score(y_test,y_pred3))
```

```
0.9835589941972921  
[[895   1]  
 [ 16 122]]  
0.991869918699187
```

```
[62]: from sklearn.linear_model import LogisticRegression  
      from sklearn.svm import SVC  
      from sklearn.naive_bayes import MultinomialNB  
      from sklearn.tree import DecisionTreeClassifier  
      from sklearn.neighbors import KNeighborsClassifier  
      from sklearn.ensemble import RandomForestClassifier  
      from sklearn.ensemble import AdaBoostClassifier  
      from sklearn.ensemble import BaggingClassifier  
      from sklearn.ensemble import ExtraTreesClassifier  
      from sklearn.ensemble import GradientBoostingClassifier  
      from xgboost import XGBClassifier
```

```
[63]: svc = SVC(kernel='sigmoid', gamma=1.0)
      knc = KNeighborsClassifier()
      mnb = MultinomialNB()
      dtc = DecisionTreeClassifier(max_depth=5)
      lrc = LogisticRegression(solver='liblinear', penalty='l1')
      rfc = RandomForestClassifier(n_estimators=50, random_state=2)
      abc = AdaBoostClassifier(n_estimators=50, random_state=2)
      bc = BaggingClassifier(n_estimators=50, random_state=2)
      etc = ExtraTreesClassifier(n_estimators=50, random_state=2)
      gbdt = GradientBoostingClassifier(n_estimators=50, random_state=2)
      xgb = XGBClassifier(n_estimators=50, random_state=2)
```

```
[64]: clfs = {
      'SVC' : svc,
      'KN' : knc,
      'NB': mnb,
      'DT': dtc,
      'LR': lrc,
      'RF': rfc,
      'AdaBoost': abc,
      'BgC': bc,
      'ETC': etc,
      'GBDT': gbdt,
      'xgb': xgb
    }
```

```
[65]: def train_classifier(clf, X_train, y_train, X_test, y_test):
      clf.fit(X_train, y_train)
      y_pred = clf.predict(X_test)
      accuracy = accuracy_score(y_test, y_pred)
      precision = precision_score(y_test, y_pred)

      return accuracy, precision
```

```
[66]: train_classifier(svc, X_train, y_train, X_test, y_test)
```

```
[66]: (0.9758220502901354, 0.9747899159663865)
```

```
[67]: accuracy_scores = []
      precision_scores = []

      for name, clf in clfs.items():

          current_accuracy, current_precision = train_classifier(clf,
↪ X_train, y_train, X_test, y_test)

          print("For ", name)
```

```

print("Accuracy - ",current_accuracy)
print("Precision - ",current_precision)

accuracy_scores.append(current_accuracy)
precision_scores.append(current_precision)

```

```

For SVC
Accuracy - 0.9758220502901354
Precision - 0.9747899159663865
For KN
Accuracy - 0.9052224371373307
Precision - 1.0
For NB
Accuracy - 0.9709864603481625
Precision - 1.0
For DT
Accuracy - 0.9274661508704062
Precision - 0.8118811881188119
For LR
Accuracy - 0.9584139264990329
Precision - 0.9702970297029703
For RF
Accuracy - 0.9758220502901354
Precision - 0.9829059829059829
For AdaBoost
Accuracy - 0.9245647969052224
Precision - 0.8488372093023255
For BgC
Accuracy - 0.9584139264990329
Precision - 0.8682170542635659
For ETC
Accuracy - 0.9748549323017408
Precision - 0.9745762711864406
For GBDT
Accuracy - 0.9468085106382979
Precision - 0.9191919191919192
For xgb
Accuracy - 0.9671179883945842
Precision - 0.9482758620689655

```

```

[68]: performance_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy':
    ↳accuracy_scores,'Precision':precision_scores}).
    ↳sort_values('Precision',ascending=False)

```

```

[69]: performance_df

```

```
[69]:
```

	Algorithm	Accuracy	Precision
1	KN	0.905222	1.000000
2	NB	0.970986	1.000000
5	RF	0.975822	0.982906
0	SVC	0.975822	0.974790
8	ETC	0.974855	0.974576
4	LR	0.958414	0.970297
10	xgb	0.967118	0.948276
9	GBDT	0.946809	0.919192
7	BgC	0.958414	0.868217
6	AdaBoost	0.924565	0.848837
3	DT	0.927466	0.811881


```
[70]: performance_df1 = pd.melt(performance_df, id_vars = "Algorithm")
```

```
[71]: performance_df1
```

```
[71]:
```

	Algorithm	variable	value
0	KN	Accuracy	0.905222
1	NB	Accuracy	0.970986
2	RF	Accuracy	0.975822
3	SVC	Accuracy	0.975822
4	ETC	Accuracy	0.974855
5	LR	Accuracy	0.958414
6	xgb	Accuracy	0.967118
7	GBDT	Accuracy	0.946809
8	BgC	Accuracy	0.958414
9	AdaBoost	Accuracy	0.924565
10	DT	Accuracy	0.927466
11	KN	Precision	1.000000
12	NB	Precision	1.000000
13	RF	Precision	0.982906
14	SVC	Precision	0.974790
15	ETC	Precision	0.974576
16	LR	Precision	0.970297
17	xgb	Precision	0.948276
18	GBDT	Precision	0.919192
19	BgC	Precision	0.868217
20	AdaBoost	Precision	0.848837
21	DT	Precision	0.811881

```
[72]: sns.catplot(x = 'Algorithm', y='value',
                  hue = 'variable',data=performance_df1, kind='bar',height=5)
plt.ylim(0.5,1.0)
plt.xticks(rotation='vertical')
plt.show()
```



```
[73]: temp_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_max_ft_3000':
    ↳ accuracy_scores,'Precision_max_ft_3000':precision_scores}).
    ↳ sort_values('Precision_max_ft_3000',ascending=False)
```

```
[74]: temp_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_scaling':
    ↳ accuracy_scores,'Precision_scaling':precision_scores}).
    ↳ sort_values('Precision_scaling',ascending=False)
```

```
[75]: new_df = performance_df.merge(temp_df,on='Algorithm')
```

```
[76]: new_df_scaled = new_df.merge(temp_df,on='Algorithm')
```

```
[77]: temp_df = pd.DataFrame({'Algorithm':clfs.keys(),'Accuracy_num_chars':
    ↳ accuracy_scores,'Precision_num_chars':precision_scores}).
    ↳ sort_values('Precision_num_chars',ascending=False)
```



```
[78]: new_df_scaled.merge(temp_df,on='Algorithm')
```

```
[78]:   Algorithm  Accuracy  Precision  Accuracy_scaling_x  Precision_scaling_x  \
0         KN  0.905222   1.000000           0.905222           1.000000
1         NB  0.970986   1.000000           0.970986           1.000000
2         RF  0.975822   0.982906           0.975822           0.982906
3         SVC 0.975822   0.974790           0.975822           0.974790
4         ETC 0.974855   0.974576           0.974855           0.974576
5         LR  0.958414   0.970297           0.958414           0.970297
6         xgb 0.967118   0.948276           0.967118           0.948276
7         GBDT 0.946809   0.919192           0.946809           0.919192
8         BgC 0.958414   0.868217           0.958414           0.868217
9   AdaBoost 0.924565   0.848837           0.924565           0.848837
10        DT  0.927466   0.811881           0.927466           0.811881
```

```
   Accuracy_scaling_y  Precision_scaling_y  Accuracy_num_chars  \
0           0.905222           1.000000           0.905222
1           0.970986           1.000000           0.970986
2           0.975822           0.982906           0.975822
3           0.975822           0.974790           0.975822
4           0.974855           0.974576           0.974855
5           0.958414           0.970297           0.958414
6           0.967118           0.948276           0.967118
7           0.946809           0.919192           0.946809
8           0.958414           0.868217           0.958414
9           0.924565           0.848837           0.924565
10          0.927466           0.811881           0.927466
```

```
   Precision_num_chars
0           1.000000
1           1.000000
2           0.982906
3           0.974790
4           0.974576
5           0.970297
6           0.948276
7           0.919192
8           0.868217
9           0.848837
10          0.811881
```

```
[79]: svc = SVC(kernel='sigmoid', gamma=1.0,probability=True)
      mnb = MultinomialNB()
      etc = ExtraTreesClassifier(n_estimators=50, random_state=2)

      from sklearn.ensemble import VotingClassifier
```

```
[80]: voting = VotingClassifier(estimators=[('svm', svc), ('nb', mnb), ('et',  
→etc)], voting='soft')
```

```
[81]: voting.fit(X_train, y_train)
```

```
[81]: VotingClassifier(estimators=[('svm',  
                                SVC(gamma=1.0, kernel='sigmoid',  
                                    probability=True)),  
                                ('nb', MultinomialNB()),  
                                ('et',  
                                ExtraTreesClassifier(n_estimators=50,  
                                                    random_state=2))],  
                    voting='soft')
```

```
[82]: y_pred = voting.predict(X_test)  
print("Accuracy", accuracy_score(y_test, y_pred))  
print("Precision", precision_score(y_test, y_pred))
```

```
Accuracy 0.9816247582205029  
Precision 0.9917355371900827
```

```
[83]: estimators=[('svm', svc), ('nb', mnb), ('et', etc)]  
final_estimator=RandomForestClassifier()
```

```
[84]: from sklearn.ensemble import StackingClassifier
```

```
[85]: clf = StackingClassifier(estimators=estimators, final_estimator=final_estimator)
```

```
[86]: clf.fit(X_train, y_train)  
y_pred = clf.predict(X_test)  
print("Accuracy", accuracy_score(y_test, y_pred))  
print("Precision", precision_score(y_test, y_pred))
```

```
Accuracy 0.9806576402321083  
Precision 0.946969696969697
```

```
[87]: import pickle  
pickle.dump(tfidf, open('vectorizer.pkl', 'wb'))  
pickle.dump(mnb, open('model.pkl', 'wb'))
```